

## Performance

cheat  
sheet!



**IRON RULE:** Execution Time = #Instructions\*CPI\*Clock Period = #Cycles\*Clock Period

- Clock Period** **#Cycles**
- Single Cycle: Latency of Slowest Instruction
  - Multi Cycle: Latency of Slowest Cycle (generally slowest datapath component)
- #Cycles Per Opcode (MULTICYCLE ONLY)**
- add/nor/sw/beq: 4 cycles
  - lw: 5 cycles
  - noop/halt: 2 cycles
  - jalr: Don't worry about it
- CPI**
- Single Cycle: 1 (it's in the name)
  - Multi Cycle: #Cycles/#Instructions

## Components used by each Opcode



Instr.	Read PC, Access Instr. Mem.	Read Reg.	ALU	Data Mem. Access	Write PC	Write Reg.
add	✓	✓	✓			✓
nor	✓	✓	✓			✓
lw	✓	✓	✓	✓ (Read)		✓
sw	✓	✓	✓	✓ (Write)		
beq	✓	✓	✓x2		✓	
jalr	✓	✓			✓	✓
noop	✓					
halt	✓					

## Datapath Benchmarking



We are using a benchmarking program that executes 10,000 instructions. 25% are **lw**, 15% **sw**, 30% **add / nor**, 20% **beq**, and 10% **noop / halt**.

Each use of the ALU costs **10ns**, reading memory costs **50ns**, writing memory costs **60ns**, register file reads and writes cost **5ns**, and everything else costs **0ns**.

**Key for today: What is the runtime of this benchmark on different systems?**

## Optimal Performance Benchmarking



**What is the optimal achievable serial runtime of this benchmark?**

Main idea: only use the wire latency for each instruction.

add / nor (50 + 5 + 10 + 5)ns = 70ns      beq (50 + 5 + 10)ns = 65ns  
lw (50 + 5 + 10 + 50 + 5)ns = 120ns      noop / halt 50ns  
sw (50 + 5 + 10 + 60)ns = 125ns

Total = 10,000 (0.25 × 120ns (lw) + 0.15 × 125ns (sw) + 0.3 × 70ns (add / nor) + 0.2 × 65ns (beq) + 0.1 × 50ns (noop / halt) )  
= 877,500ns

## Pipeline Performance



#Cycles = #instr + #pipeline\_stages - 1  
(for the 5-stage pipeline) -> #Cycles = #instr + 4

Base CPI = #Cycles / #Instr ≈ 1

Then add 3 S's: Startup cycles, Stalls, Squashes

EX: For a program with 1000 instructions with no hazards, we would need **1004 cycles** to execute and the CPI would be **1004/1000 = 1.004**

## We Can Guess If a Branch Is Taken Using Predictors



- In order to predict a branch, we need to infer three things:
- Whether the instruction is a branch (*We haven't Decoded yet!*)
  - If the branch instruction will be taken or not
  - Where the branch will jump to if it is taken
- If we can't figure out some of this information, we cannot make a prediction, often revert to assuming it will not be taken
- For example, if we don't know where a branch will jump to, we can never predict it as taken, since we won't know where to fetch from next

## Static Branch Predictors Don't Change Prediction State




- Always Not Taken:** Predicts that a branch will never be taken (Project 3)
- Always Taken:** Predicts that a branch will always be taken
- Backwards Taken, Forwards Not Taken:** Predicts a branch will be taken or not based on if it branches forwards or backwards

These predictions can all be done at compile time

## Dynamic Predictors Take Advantage of the Past



- Last Time Predictor:** Predict based on whatever the branch did last time
- Problem:* Could cause 0% accuracy for "T N T N T N T ..."



# FSM For a 2-bit Branch Predictor

