

EECS 370 - Lecture 5

C to Assembly



M Live Poll + Q&A: slido.com #eeecs370

Poll and Q&A Link

M

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks**
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout

Agenda

- Memory alignment**
 - Aligning Structs
- Control flow instructions
 - C-code examples
- Extra Problems

Calculating Load/Store Addresses for Variables

Datatype	size (bytes)
char	1
short	2
int	4
double	8

```
short a[100];
char b;
int c;
double d;
short e;
struct {
    char f;
    int g[1];
    char h;
} i;
```

- Problem:** Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed
 - a = 2 bytes * 100 = 200
 - b = 1 byte
 - c = 4 bytes
 - d = 8 bytes
 - e = 2 bytes
 - i = 1 + 4 + 1 = 6 bytes
 - total = 221, right or wrong? X

4

M

5

Memory layout of variables

- Compilers don't like variables placed in memory arbitrarily
- As we'll see later in the course, memory is divided into fixed sized chunks
 - When we load from a particular chunk, we really read the whole chunk
 - Usually an integer number of words (32 bits)
- If we read a single char (1 byte), it doesn't matter where it's placed

0x1000	0x1001	0x1002	0x1003
'a'	'b'	'c'	'd'

ldurb [0x0, 0x1002]

• Reads [0x1000-0x1003], then throws away all but 0x1002, **fine**

Memory layout of variables

- BUT, if we read a 32-bit integer word, and that word starts at 0x1002:

0x1000	0x1001	0x1002	0x1003
0xFF	0xFF	0x12	0x34
0x1004	0x1005	0x1006	0x1007
0x56	0x78	0xFF	0xFF

- First we need to read [0x1000-0x1003], throw away 0x1000 and 0x1001, **then** read [0x1004-0x1007]
- Need to read from memory twice! Slow! Complicated! **Bad!**

6

M

7

Solution: Memory Alignment

Poll: Where can chars start?

ex: int 4B

can be placed

on 0x1000

on 0x1002

on 0x1004

on 0x1006

on 0x1008

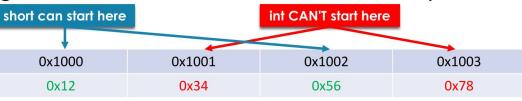
- Most modern ISAs require that data be aligned

~~• An N-byte variable must start at an address A, such that $(A \% N) == 0$~~

~~• For example, starting address of a 32 bit int must be divisible by 4~~



- Starting address of a 16 bit short must be divisible by 2



Golden Rule of Alignment

(Unit 6)

- Every (primitive) object starts at an address divisible by its size**
- "Padding" is placed in between objects if needed

```
char c;
short s;
int i;
```

0x1000	0x1001	0x1002	0x1003	0x1004	0x1005	0x1006	0x1007
[c]	[padding]	[s]			[i]		

- But what about non-primitive data types?

- Arrays? Treat as independent objects
- Structs? Trickier... ?

8

M

9

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - C-code examples
- Extra Problems

Problem with Structs

`char c;`

```
struct {
    char c;
    int i;
} s[2];
```

- If we align each element of a struct according to the Golden Rule, we can still run into issues
 - E.g.: An array of structs

Amount of padding is different across different instances

1000 1001 1002 1003 1004 1005 1006 1007 1008 1009 100A 100B 100C 100D 100E 100F
c s[0].c [pad] s[0].i s[1].c [pad] s[1].i

- Why is this bad?
- It makes "for" loops very difficult to write!
 - Offsets need to be different on each iteration

M

10

M

11

Structure Alignment

- Solution: in addition to laying out each field according to Golden Rule...
 - Identify largest (primitive) field
 - Starting address of overall struct is aligned based on the largest field
 - Padded in the back so total size is a multiple of the largest primitive

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c [pad]	[pad]	[pad]	[pad]	s[0].c [pad]	[pad]	[pad]	[pad]	[pad]	s[0].i	s[1].c [pad]	[pad]	[pad]	[pad]	s[1].i	

Guaranteed to lay out each instance identically

```
char c;
struct {
    char c;
    int i;
} s[2];
```

12

Structure Example

```
struct {
    char w; 1
    int x[3]; 4x3
    char y; 1
    short z; 2
}
```

Poll: What boundary should this struct be aligned to?

- 1 byte
- 4 bytes
- 12 bytes
- 2 bytes
- 19 bytes

$$\frac{1+3+4\times 3+1+2+1}{4} = 20$$

- Assume struct starts at location 1000,

- char w → 1000
- x[0] → 1004–1007, x[1] → 1008–1011, x[2] → 1012–1015
- char y → 1016
- short z → 1018–1019

Total size = 20 bytes!

M Live Poll + Q&A: [slido.com #eecs370](https://slido.com/#eecs370)

13

Calculating Load/Store Addresses for Variables

Datatype	size (bytes)
char	1
short	2
int	4
double	8

```
short a[100]; 200 (~298)
char b; 1 (300~302)
int c; 4 (304~308) since
double d; 8 (312~319) 4(310)
short e; 2 (320~321)
struct {
    char f; 1 (13)
    int g[1]; 4 (324-328)
    char h; 1 (14)
} i;

```

• Problem: Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

a = 200 bytes (100-299)
b = 1 byte (300-300)
c = 4 bytes (304-307)
d = 8 bytes (312-319)
e = 2 bytes (320-321)
struct: largest field is 4 bytes, start at 324
f = 1 byte (324-324)
g = 4 bytes (328-331)
h = 1 byte (332-332)
i = 12 bytes (324-335)

236 bytes total!! (compared to 221, originally)

14

M

15

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - C-code examples
- Extra Problems

ARM/LEGv8 Sequencing Instructions

- Sequencing instructions change the flow of instructions that are executed
 - This is achieved by modifying the program counter (PC)
- Unconditional branches are the most straightforward they ALWAYS change the PC and thus "jump" to another instruction out of the usual sequence
- Conditional branches

```
If (condition_test) goto target_address
condition_test examines the four flags from the processor status word (SPSR)
target_address is a 19 bit signed word displacement on current PC
```

M

16

M

17

LEGv8 Conditional Instructions

- Two varieties of conditional branches

Conditional branch	compare and branch CBZ X1, 25 on equal 0	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch CBNZ X1, 25 on not equal 0	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally B.cond 25	if (condition true) go to test condition codes; if true, branch PC + 100	

- Let's look at the first type: CBZ and CBNZ

- CBZ: Conditional Branch if Zero
- CBNZ: Conditional Branch if Not Zero

18



19

LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again: ADDI X3, X3, #-1

CBNZ X3, Again

Poll: What's the offset?

- a) -16
b) -12
c) -4
d) -3
e) 0

M Live Poll + Q&A: slido.com #eeecs370

20

M Live Poll + Q&A: slido.com #eeecs370

21

LEGv8 Conditional Instructions

- Motivation:
 - Some types of branches makes sense to check if a certain value is zero or not
 - while(a)
 - But not all:
 - if(a > b)
 - if(a == b)
- Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior

M 22



ARM Condition Codes Determine Direction of Branch

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP / CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
 - N – set if the result is **negative** (i.e., bit 63 is non-zero)
 - Z – set if the result is **zero** (i.e., all 64 bits are zero)
 - C – set if last addition/subtraction had a **carry/borrow** out of bit 63
 - V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)
- (Don't worry about the C and V for this class)

24



25

LEGv8 Conditional Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
- The relative address is a 19 bit signed integer—the number of instructions. Recall instructions are 32 bits of 4 bytes

Conditional branch	compare and branch CBZ X1, 25 on equal 0	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch CBNZ X1, 25 on not equal 0	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally B.cond 25	if (condition true) go to test condition codes; if true, branch PC + 100	

- Example: CBNZ X3, Again

- If X3 doesn't equal 0, then branch to label "Again"
- "Again" is an offset from the PC of the current instruction (CBNZ)
- Why does "25" in the above table result in PC + 100?

LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again: ADDI X3, X3, #-1

CBNZ X3, Again (E-3)

- Answer = -3

- The offset field is 19 bits signed so the bit pattern would be 111 1111 1111 1111 1101
- Two 0's are appended to the above 19 bits and then the result would be sign-extended (with one's) to 64 bits and added to the value of PC at CBNZ
- Why the two 00's?

LEGv8 Conditional Instructions Using FLAGS

- FLAGS: NZVC record the results of (arithmetic) operations Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	A00 X1, X2, X3	X1 - X2 + X3	Three register operands
	subtract	SUB X1, X2, X3	X1 - X2 - X3	Three register operands
	add immediate	A001 X1, X2, 20	X1 - X2 + 20	Used to add constants
	subtract immediate	SUB1 X1, X2, 20	X1 - X2 - 20	Used to subtract constants
	add and set flags	A005 X1, X2, X3	X1 - X2 + X3	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	X1 - X2 - X3	Subtract, set condition codes
	add immediate and set flags	A005S X1, X2, 20	X1 - X2 + 20	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	X1 - X2 - 20	Subtract constant, set condition codes

ARM Condition Codes Determine Direction of Branch--continued

Encoding	Name (& alias)	Meaning (integer)	Flags
0000	EQ	Equal	Z==1
0001	NE	Not equal	Z==0
0010	HS	Unsigned higher or same	C==1
0011	LS	(Carry set)	
0100	LO	Unsigned lower (Carry clear)	C==0
0101	MI	Minus (negative)	N==1
0110	PL	Plus (positive or zero)	N==0
0111	VS	Overflow set	V==1
1000	HI	Overflow clear	V==0
1001	LS	Unsigned lower or same	C==1 & Z==0
1010	GE	Signed greater than or equal	N==V
1011	LT	Signed less than	N!=V
1100	GT	Signed greater than	Z==0 & N==V
1101	LE	Signed less than or equal	! (Z==0 & N==V)
1110	AL	Always	
1111	NV		

Need to know the 7 with the red arrows

CMP X1, X2
B.LE Label1

For this example, we branch if X1 is >= to X2

Conditional Branches: How to use

- CMP instruction lets you compare two registers.
 - Could also use SUBS etc.
 - That could save you an instruction.
- B.cond lets you branch based on that comparison.

store \Rightarrow CPSR

- Example:

```
CMP X1, X2
B.GT Label1
```

- Branches to Label1 if X1 is greater than X2.

M

26

Branch—Example

- Convert the following C code into LEGV8 assembly (assume x is in X1, y in X2):

```
int x, y;
if (x == y)
    x++;
else
    y++;
// ...
L1: ADD X2, X2, #1
L2: ...
```

Using Labels

Note that conditions in assembly are often the inverse of the "if" condition. Why?

	Without Labels
CMP	X1, X2
B.NE	3
ADD	X1, X1, #1
B	2
ADD	X2, X2, #1

Assemblers must deal with labels and assign displacements

M

29

Loop—Example

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}

# of branch instructions
= 3*10 + 1 = 31
a.k.a. while-do template
```

```
MOV X1, XZR
MOV X2, XZR (d)
Loop1: CMPI i X1, #10
        B.EQ endLoop
        LSL X6, X1, #3
        LDUR X5, [X6, #100] X6 = i * 8
        CMPI X5, #0
        B.LT endif
        ADD X2, X2, X5
        ADDI X1, X1, #1
        B Loop1
endif:
endLoop:
```

M

30

Extra Example: Do-while Loop

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}

# of branch instructions
= 2*10 = 20
a.k.a. do-while template
```

	MOV X1, XZR	MOV X2, XZR
Loop1:	LSL X6, X1, #3	LDUR X5, [X6, #100]
	CMPI X5, #0	B.LT endif
	ADD X2, X2, X5	ADDI X1, X1, #1
	CMPI X1, #10	B.LT Loop1

32

Extra Example: Do-while Loop

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}

# of branch instructions
= 2*10 = 20
a.k.a. do-while template
```

M

33

Extra Problem

- Write the ARM assembly code to implement the following C code:

```
/* address, i.e. phr */
ptr ? null cmp r1, #0
beq Endif
ldursw r2, [r1, #0]
cmp r2, #0
b.le Endif
add r2, r2, #1
str r2, [r1, #0]
Endif : ....
```

35

Extra Class Problem

- How much memory is required for the following data, assuming that the data starts at address 200 and is a 32 bit address space?

int a; (4B (200~203))
struct {double b, char c, int d} e; = 16 (208~223)
char* f; (4B (224~227))
short g[20]; (40B (228~247))

$\Rightarrow 68B$

Poll: How much memory?

- x < 40 bytes
- 40 < x < 50 bytes
- 50 < x < 60 bytes
- 60 < x bytes

Next Time

- More C-to-Assembly
 - Function calls
- Lingering questions / feedback? Post it to Slido

M

36

M

Live Poll + Q&A: slido.com #eecs370

37