

EECS 370 - Lecture 14

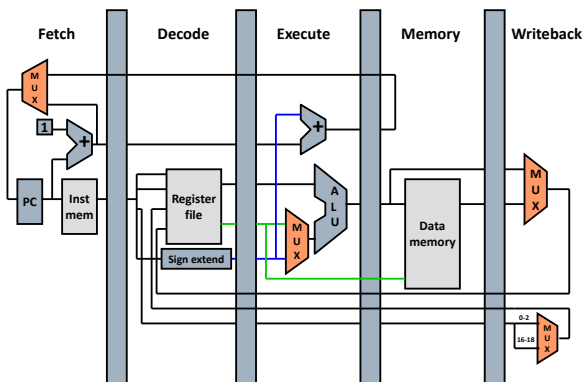
Pipelining and Data Hazards II



Review: Pipelining

- Goal:
 - Achieve low clock period of multi-cycle processor...
 - ... while maintaining low cycles-per-instruction (CPI) of single cycle processor (close to 1)
 - Achieve this by overlapping execution of multiple instructions simultaneously

Review: Our new pipelined datapath



4


Review: Sample Code (Simple)

Let's run the following code on pipelined LC2K:

- add 1 2 3 ; reg 3 = reg 1 + reg 2
- nor 4 5 6 ; reg 6 = reg 4 nor reg 5
- lw 2 4 20 ; reg 4 = Mem[reg2+20]
- add 2 5 5 ; reg 5 = reg 2 + reg 5
- sw 3 7 10 ; Mem[reg3+10] = reg 7

Data Hazards

Recall: registers are read/sourced in the "decode" stage



If not careful, nor will read a stale value of register 3

Three approaches to handling data hazards

- Avoid
 - Make sure there are no hazards in the code
- Detect and Stall
 - If hazards exist, stall the processor until they go away.
- Detect and Forward
 - If hazards exist, fix up the pipeline to get the correct value (if possible)

Problems with detect and stall

- CPI increases every time a hazard is detected!
- Is that necessary? Not always!
 - Re-route the result of the **add** to the **nor**
 - **nor** no longer needs to read R3 from reg file
 - It can get the data later (when it is ready)
 - This lets us complete the decode this cycle
 - But we need more control logic

Handling data hazards III: Detect and forward

- Detect: same as detect and stall
 - Except that all 4 hazards have to be treated differently
 - i.e., you can't logical-OR the 4 hazard signals
- Forward:
 - New **bypass datapaths** route computed data to where it is needed
 - New MUX and control to pick the right data
- **Beware:** Stalling may still be required even in the presence of forwarding

Forwarding example

- We will use this program for the next example (same as last pipeline diagram example)

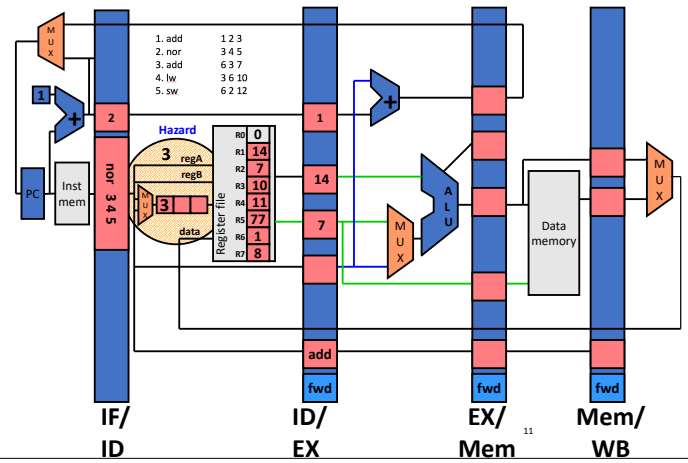
```

1. add 123
2. nor 345
3. add 637
4. lw 3610
5. sw 6212
    
```

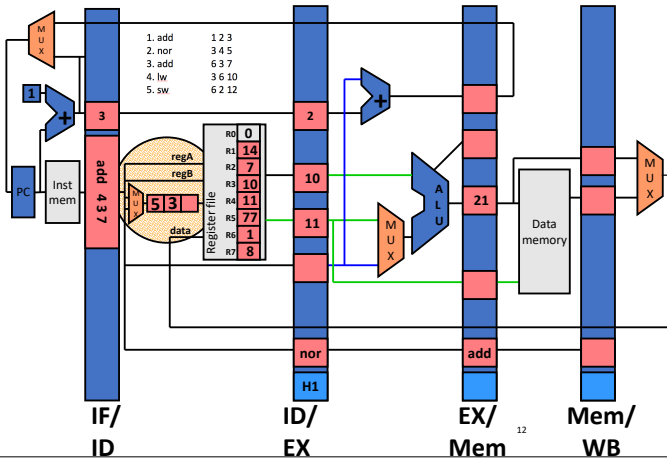


10

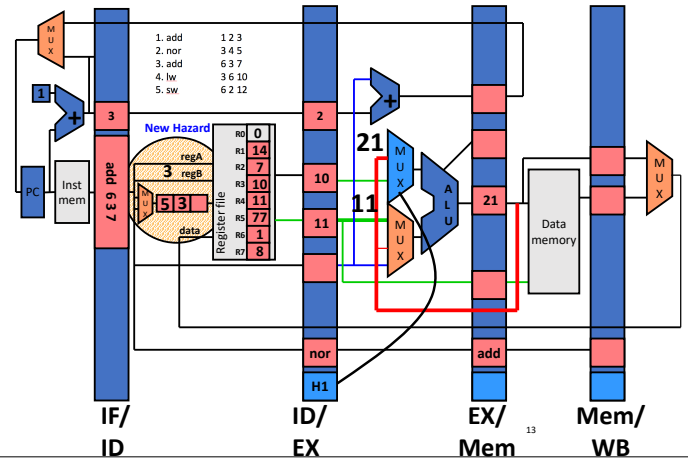
First half of cycle 3



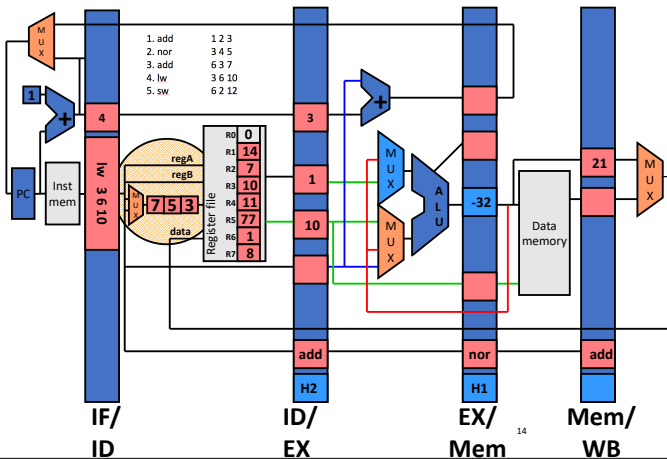
End of cycle 3



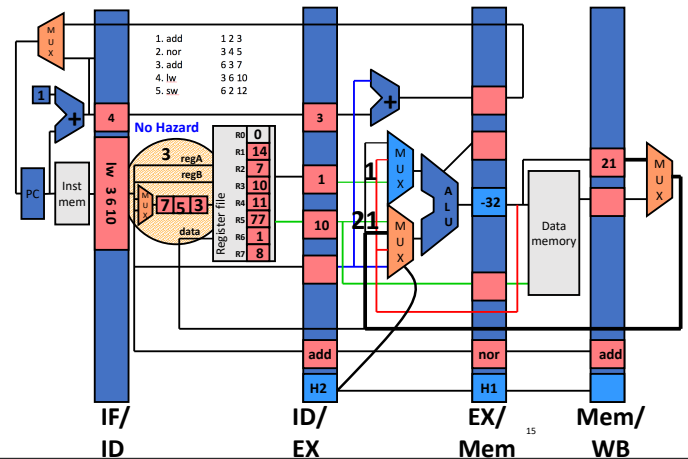
First half of cycle 4



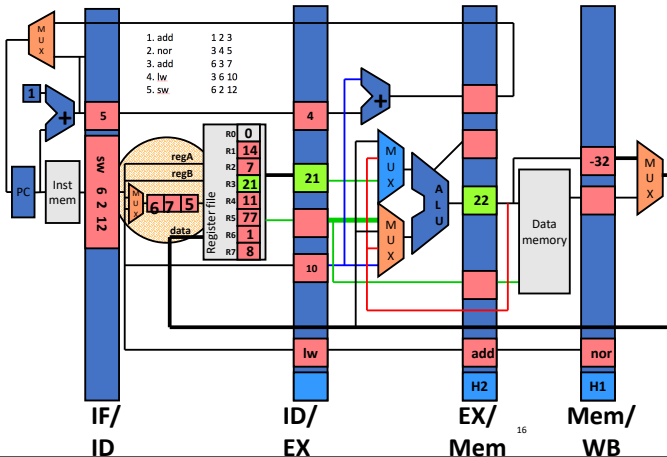
End of cycle 4



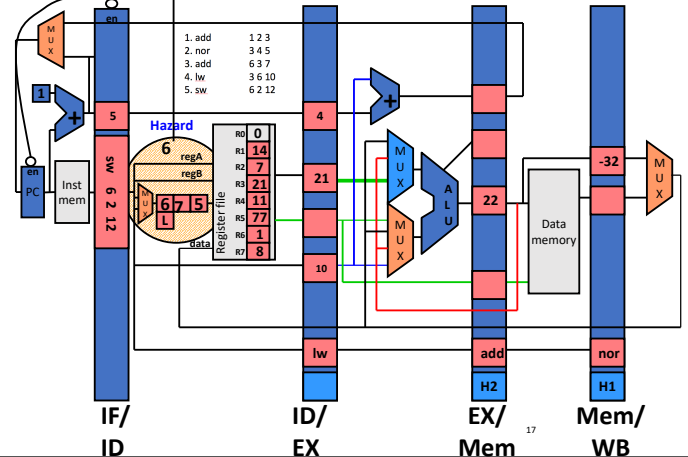
First half of cycle 5



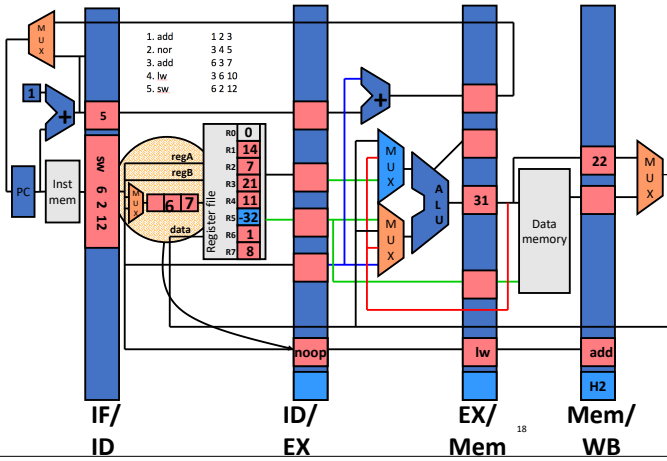
End of cycle 5



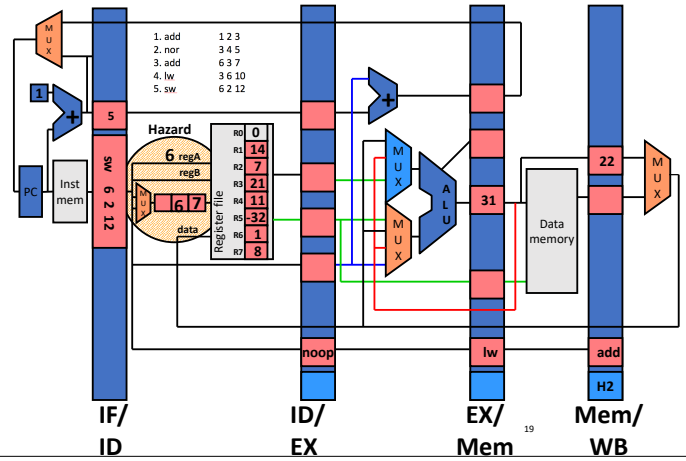
First half of cycle 6



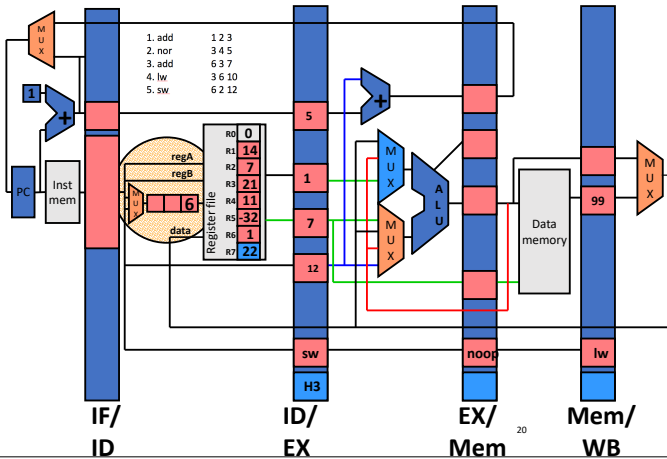
End of cycle 6



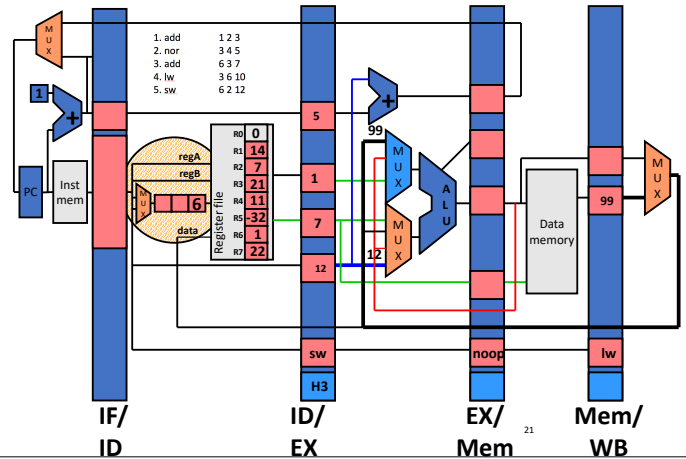
First half of cycle 7



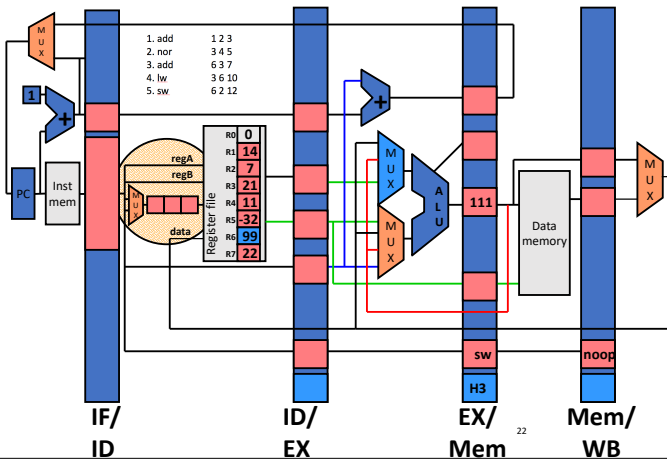
End of cycle 7



First half of cycle 8



End of cycle 8




Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nor 3 4 5		IF	ID	EX	ME	WB							
add 6 3 7			IF	ID	EX	ME	WB						
lw 3 6 10				IF	ID	EX	ME	WB					
sw 6 2 12					IF	ID*	ID	EX	ME	WB			

Other issues

- What other instruction(s) have we been ignoring so far??
- Branches!! (Let's not worry about jumps yet)
- Sequence for BEQ:
 - Fetch: read instruction from memory
 - Decode: read source operands from registers
 - Execute: calculate target address and test for equality
 - Memory: Send target to PC if test is equal
 - Writeback: nothing
 - Branch Outcomes
 - Not Taken
 - PC = PC + 1
 - Taken
 - PC = Branch Target Address

Control Hazards



Approaches to handling control hazards

- 3 strategies – similar to handling data hazards

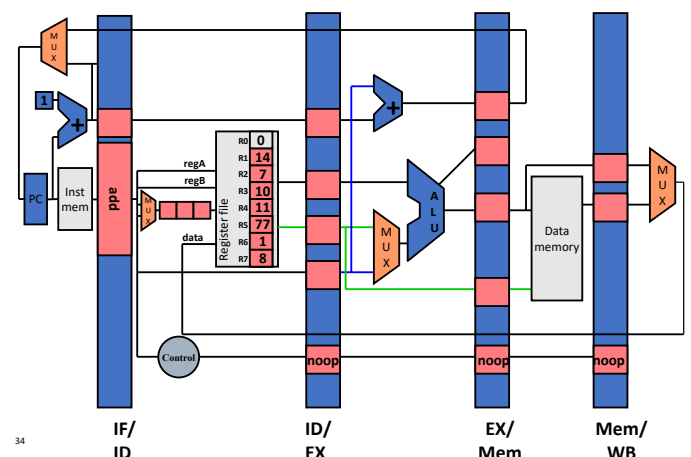
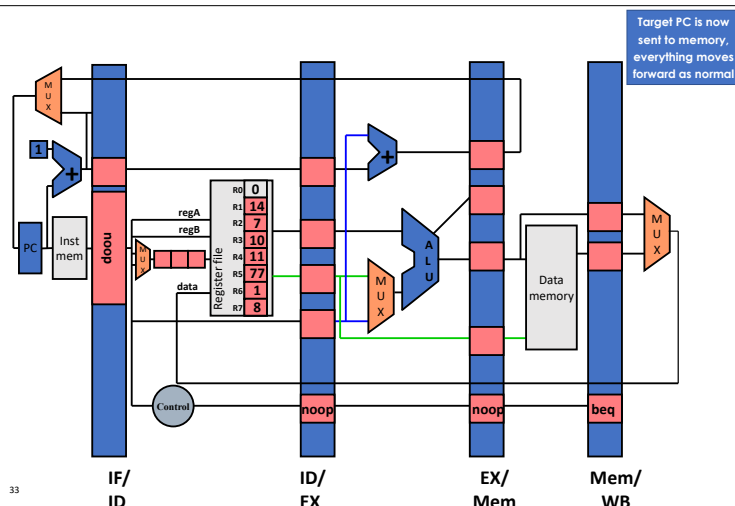
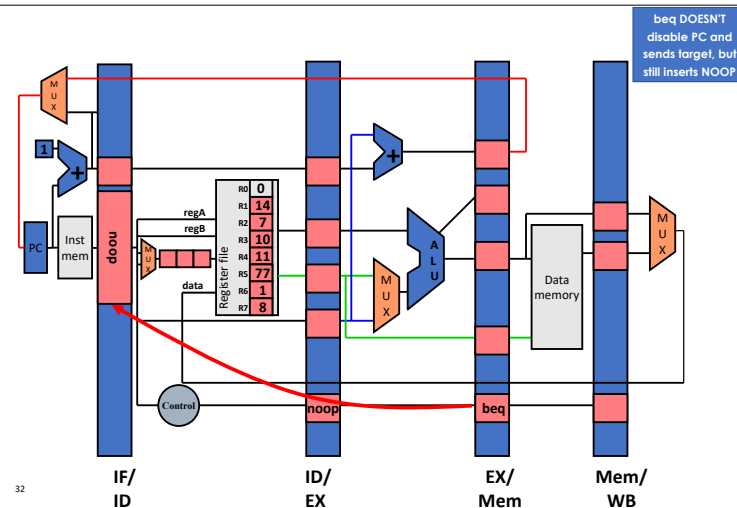
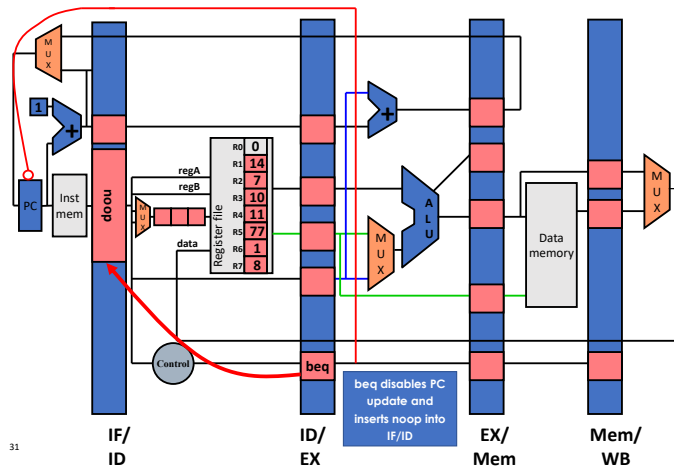
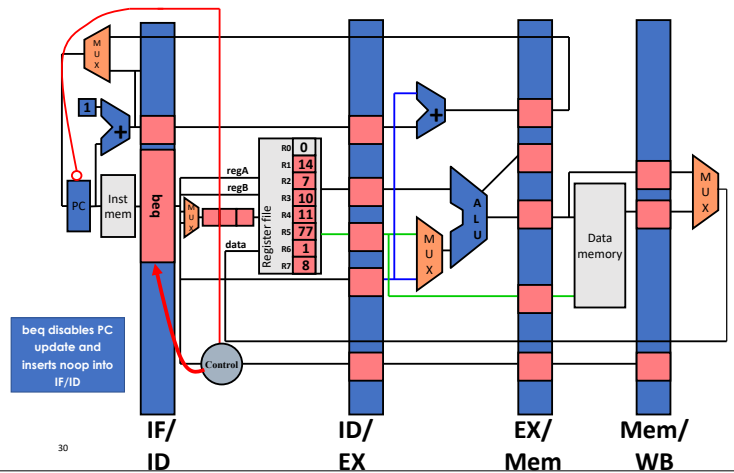
1. Avoid
 - Make sure there are no hazards in code
2. Detect and stall
 - Delay fetch until branch resolved
3. Speculate and squash-if-wrong
 - Guess outcome of branch
 - Fetch instructions assuming we're right
 - Stop them if they shouldn't have been executed

Avoiding Control Hazards


- Don't have branch instructions!
 - Possible, but not practical
- ARM offers **predicated** instructions (instructions that throw away result if some condition is not met)
 - Allows replacement of if/else conditions
 - Hard to use for everything
 - Not covered more in this class

Detect and Stall

- Detection
 - Wait until decode
 - Check if opcode == beq or jalr
- Stall
 - Keep current instruction in fetch
 - Insert noops
 - Pass noop to decode stage, not execute!



Control Hazards



35

Problems with Detect and Stall

- CPI increases every time a branch is detected!
- Is that necessary? Not always!
 - Branch not always taken
 - Let's assume it is NOT taken...
 - In this case, we can ignore the beq (treat it like a noop)
 - Keep fetching PC + 1
 - What if we're wrong?
 - OK, as long as we do not COMPLETE any instruction we mistakenly execute
 - I.e. DON'T write values to register file or memory

36

Agenda

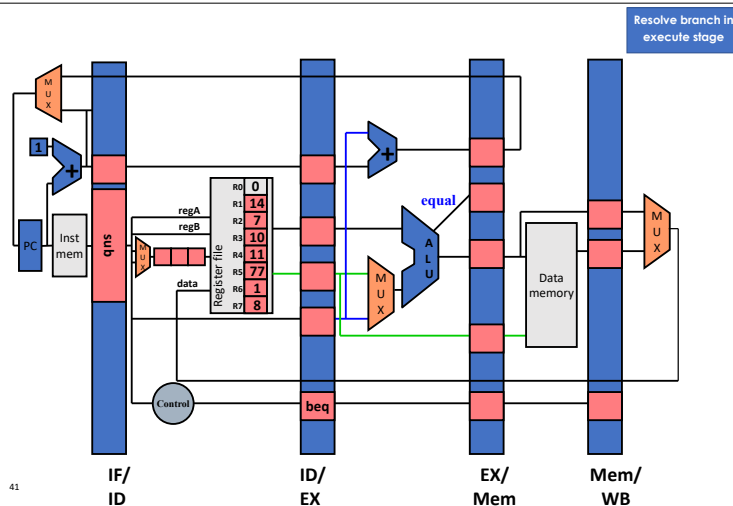
- Control Hazards and Basic Approaches
- Detect-and-Stall
- **Speculate-and-Squash**
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor

39

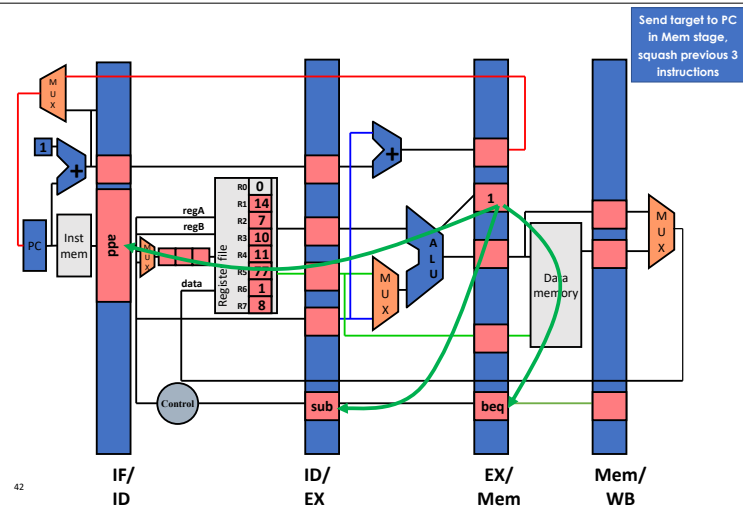
Speculate and Squash

- Speculate: assume not equal
 - Keep fetching from PC+1 until we know that the branch is really taken
- Squash: stop bad instructions if taken
 - Send a noop to Decode, Execute, and Memory
 - Sent target address to PC

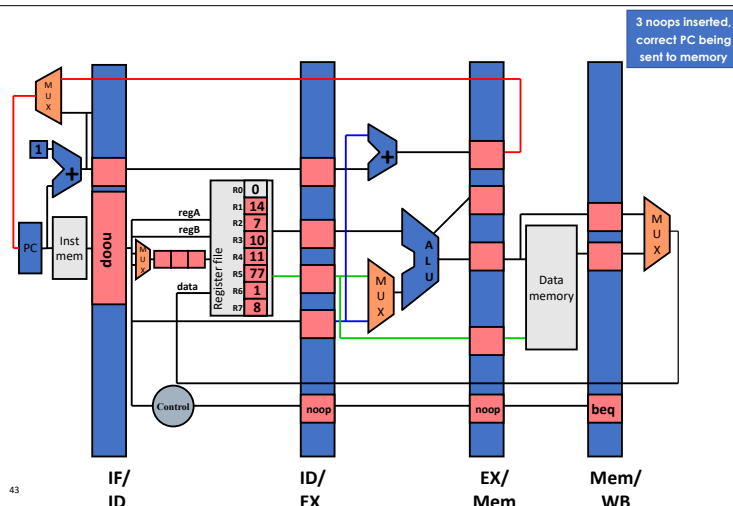
40



41



42



43

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- **Exceptions**
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor

44

What can go wrong?

- ❑ **Data hazards:** since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read the wrong value if is about to be written.
- ❑ **Control hazards:** A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- ❑ **Exceptions:** How do you handle exceptions in a pipelined processor with 5 instructions in flight?



45

Exceptions

- ❑ Exception: when something unexpected happens during program execution.
 - Example: divide by zero.
 - The situation is more complex than the hardware can handle
 - So the hardware branches to a function, an “exception handler” which is code to try to deal with the problem.
- ❑ The exact way to set up such an exception handler will vary by ISA.
 - With C on x86 you would use <signal.h> functions to handle the “SIGFPE” signal.
 - There is a pretty good [Hackaday article](#) on this if you want to learn more.



46

Exceptions and Pipelining

- ❑ The hardware branches to the “exception handler”
 - This means that any instruction which can “throw” an exception *could* be a branch.
 - Throwing an exception should be rare (“exceptional”)
- ❑ So we would treat it much like a branch we predicted as “not taken”
 - Squash instructions behind it and then branch
 - It will introduce stalls, but since it should be rare, we don’t worry about it.
 - “Make the common case fast”.



47

Classic performance problem

- ❑ Program with following instruction breakdown:

lw	10%
sw	15%
beq	25%
R-type	50%
 - ❑ Speculate “always not-taken” and squash. 80% of branches not-taken
 - ❑ Full forwarding to execute stage. 20% of loads stall for 1 cycle
 - ❑ What is the CPI of the program?
 - ❑ What is the total execution time per instruction if clock frequency is 100MHz?
- $$\text{CPI} = 1 + 0.10 (\text{loads}) * 0.20 (\text{load use stall}) * 1 + 0.25 (\text{branch}) * 0.20 (\text{miss rate}) * 3$$
$$\text{CPI} = 1 + 0.02 + 0.15 = 1.17$$
$$\text{Time} = 1.17 * 10\text{ns} = 11.7\text{ns per instruction}$$



51

Classic performance problem (cont.)

- ❑ Assume branches are resolved at Execute?
 - What is the CPI?
 - What happens to cycle time?

$$\text{CPI} = 1 + 0.10 (\text{loads}) * 0.20 (\text{load use stall}) * 1 + 0.25 (\text{branch}) * 0.20 (\text{miss rate}) * 2$$
$$\text{CPI} = 1 + 0.02 + 0.1 = 1.12$$



54

Performance with deeper pipelines

- ❑ Assume the setup of the previous problem.
 - ❑ What if we have a 10 stage pipeline?
 - Instructions are fetched at stage 1.
 - Register file is read at stage 3.
 - Execution begins at stage 5.
 - Branches are resolved at stage 7.
 - Memory access is complete in stage 9.
 - ❑ What’s the CPI of the program?
 - ❑ If the clock rate was doubled by doubling the pipeline depth, is performance also doubled?
- $$\text{CPI} = 1 + 0.10 (\text{loads}) * 0.20 (\text{load use stall}) * 4 + 0.25 (\text{branch}) * 0.20 (\text{N stalls}) * 6$$
$$\text{CPI} = 1 + 0.08 + 0.30 = 1.38$$
$$\text{Time} = 1.38 * 5\text{ns} = 6.9\text{ ns per instruction}$$



57