

EECS 370

Control Hazards and Performance

Classic performance problem

- ❑ Program with following instruction breakdown:

lw	10%
sw	15%
beq	25%
R-type	50%
- ❑ Speculate “always not-taken” and squash. 80% of branches not-taken
- ❑ Full forwarding to execute stage. 20% of loads stall for 1 cycle
- ❑ What is the CPI of the program?
- ❑ What is the total execution time per instruction if clock frequency is 100MHz?

$$\text{CPI} = 1 + 0.10 (\text{loads}) * 0.20 (\text{load use stall}) * 1 + 0.25 (\text{branch}) * 0.20 (\text{miss rate}) * 3$$

$$\text{CPI} = 1 + 0.02 + 0.15 = 1.17$$

$$\text{Time} = 1.17 * 10\text{ns} = 11.7\text{ns per instruction}$$

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - **Problem 2**
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor

Classic performance problem (cont.)

- ❑ Assume branches are resolved at Execute?
 - What is the CPI?
 - What happens to cycle time?

$$\text{CPI} = 1 + 0.10 (\text{loads}) * 0.20 (\text{load use stall}) * 1 + 0.25 (\text{branch}) * 0.20 (\text{miss rate}) * 2$$

$$\text{CPI} = 1 + 0.02 + 0.1 = 1.12$$

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - **Problem 3**
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- Improving Direction Predictor

Performance with deeper pipelines

- ❑ Assume the setup of the previous problem.
- ❑ What if we have a 10 stage pipeline?
 - Instructions are fetched at stage 1.
 - Register file is read at stage 3.
 - Execution begins at stage 5.
 - Branches are resolved at stage 7.
 - Memory access is complete in stage 9.
- ❑ What's the CPI of the program?
- ❑ If the clock rate was doubled by doubling the pipeline depth, is performance also doubled?

$$\text{CPI} = 1 + 0.10 (\text{loads}) * 0.20 (\text{load use stall}) * 4 + 0.25 (\text{branch}) * 0.20 (\text{N stalls}) * 6$$

$$\text{CPI} = 1 + 0.08 + 0.30 = 1.38$$

$$\text{Time} = 1.38 * 5\text{ns} = 6.9\text{ns per instruction}$$

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- **Improving Performance with Branch Predicting**
- Simple Direction Predictor
- Improving Direction Predictor

Can We Improve Branch Performance?

- CPI increases every time a branch is taken!
 - About 50%-66% of time
- Is that necessary?
- **No!** We can try to predict when branch is taken
 - But we would need to send target PC to memory before decoding branch
 - How do we:
 1. Know an instruction is a branch before decoding?
 2. Reliably guess whether it should be taken?
 3. Figure out the target PC before executing the branch?

Poll: If you had to guess, in real programs, what's the ratio of taken to not-taken branches?

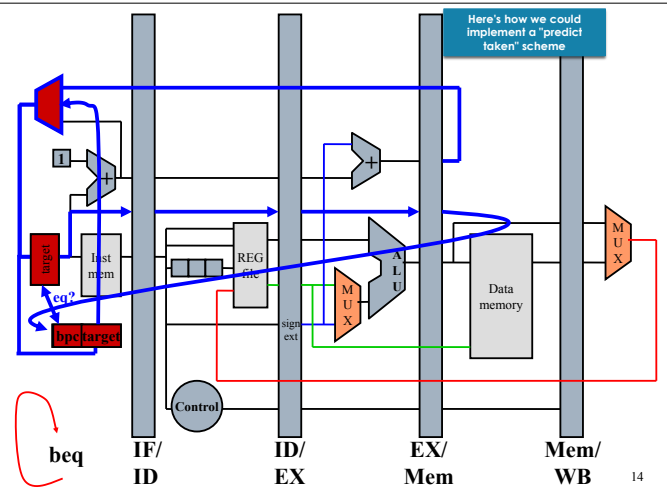
- a) Very rarely taken
- b) Slightly biased towards not taken
- c) Slightly biased towards taken
- d) Very rarely not taken

Sometimes predict taken?

- When fetching an instruction, need to predict 3 things:
 - Whether the fetched instruction is a branch
 - Branch direction (if conditional)
 - Branch target address (if direction is taken)
- Observation: Target address remains the same for conditional branch across multiple executions
 - Idea: store the target address of branch once we execute it, along with PC of instruction
 - Called Branch Target Buffer (BTB)

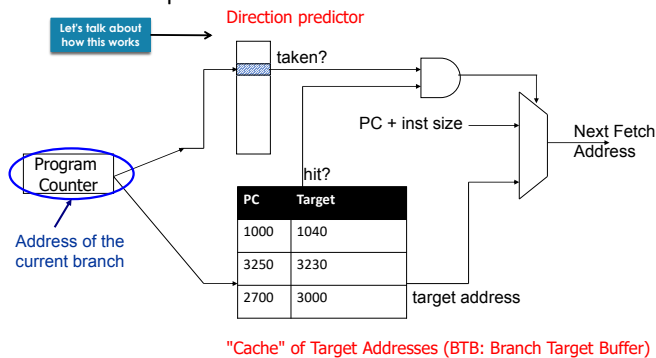


13



14

Sometimes predict taken?



15

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor**
- Improving Direction Predictor

16

Branch Direction Prediction

Poll: Overall, which do you expect to be more accurate?

- "Branch direction" refers to whether the branch was taken or not
- Two methods for predicting direction:
 - Static - We predict once during compilation, and that prediction never changes
 - Dynamic - We predict (potentially) many times during execution, and the prediction may change over time
- Static vs dynamic strategies are a very common topic in computer architecture

Branch Direction Prediction (Static)


- Always not-taken**
 - Simple to implement: no need for BTB, no direction prediction
 - Low accuracy: ~30-40%
 - Compiler can layout code such that the likely path is the "not-taken" path
- Always taken**
 - No direction prediction
 - Better accuracy: ~60-70%
 - Backward branches (i.e. loop branches) are usually taken
 - Backward branch: target address lower than branch PC
- Backward taken, forward not taken (BTFN)**
 - Predict backward (loop) branches as taken, others not-taken

18

Branch Direction Prediction (Dynamic)

- Last time predictor**
 - Single bit per branch (stored in BTB)
 - Indicates which direction branch went last time it executed
TTTTTTTTTTNNNNNNNNNN → 90% accuracy
- Always mispredicts the last iteration and the first iteration of a loop branch
 - Accuracy for a loop with N iterations = $(N-2)/N$
- + Loop branches for loops with large number of iterations
- Loop branches for loops with small number of iterations
TNTNTNTNTNTNTNTNTNTN → 0% accuracy

State Machine for Last-Time Prediction



20

19

Agenda

- Control Hazards and Basic Approaches
- Detect-and-Stall
- Speculate-and-Squash
- Exceptions
- Practice Performance Problems
 - Problem 1
 - Problem 2
 - Problem 3
- Improving Performance with Branch Predicting
- Simple Direction Predictor
- **Improving Direction Predictor**



21

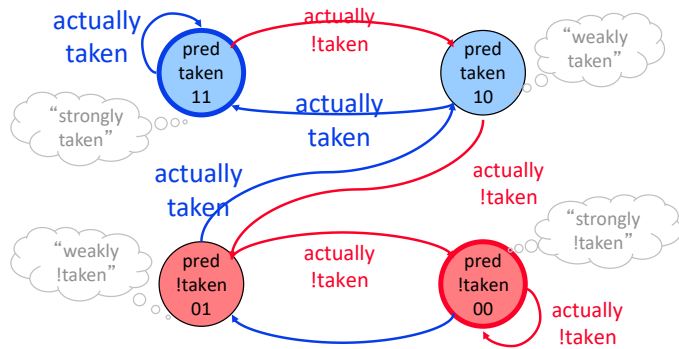
Improving the Last Time Predictor

- Problem: A last-time predictor changes its prediction from T→NT or NT→T too quickly
 - Even though the branch may be mostly taken or mostly not taken
- Solution Idea: Add hysteresis to the predictor so that prediction does not change on a single different outcome
 - Use two bits to track the history of predictions for a branch instead of a single bit
 - Can have 2 states for T or NT instead of 1 state for each



22

State Machine for 2-bit Saturating Counter



23

Two-Bit Counter Based Prediction

Poll: How many branches do we get wrong?

- What's the prediction accuracy of a branch with the following sequence of taken/not taken outcomes:

• T T T N T T N N N T N T N N

Br	T	T	T	T	N	T	T	N	N	N	T	N	T	N	N
State	10	11	11	11	X	10	11	X	X	01	X	01	X	01	00
Pred	T	T	T	T	T	T	T	T	T	N	N	N	N	N	N



Live Poll + Q&A: [slido.com #eecs370](https://slido.com/#eecs370)

24

Can We Do Better?

- Absolutely... take 470
 - Tons of sophisticated branch predictor designs
- I've worked on a few that found their way into some Chromebooks!



25

Branch Prediction

- Predict not taken: ~50% accurate
- Predict backward taken: ~65% accurate
- Predict same as last time: ~80% accurate
- Realistic designs: ~96% accurate



26

Remember this Example from Lecture 1?

- We know understand why sorting improves the inner-loop so much
 - The branch predictor is better at guessing what's gonna happen when data is sorted!

```
for (unsigned c = 0; c < arraySize; ++c)
    data[c] = std::rand() % 256;
std::sort(data, data + arraySize);

// Test
clock_t start = clock();
long long sum = 0;
// Primary loop
for (unsigned c = 0; c < arraySize; ++c)
{
    if (data[c] >= 128)
        sum += data[c];
}

double elapsedTime =
    static_cast<double>(clock() - start);
```



27

Advanced Pipelining

Not on the exam.

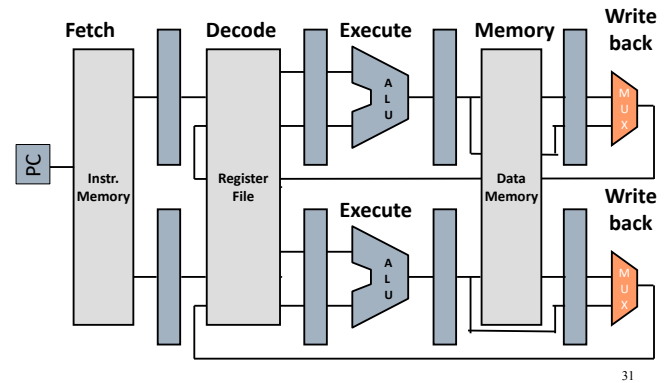
29

Creating more pipelines

- ❑ Instruction Level Parallelism – Superscalar Pipeline
 - Have two or more pipelines in same processor
 - pipelines need to work in tandem to improve single program performance
- ❑ Thread Level Parallelism – Multi-core
 - Have two or more processors (Independent Pipelines)
 - Need more programs or a parallel program
 - does not improve single program performance
- ❑ Data Level Parallelism – Single Inst. Multiple Data (SIMD)
 - Have two or more execution pipelines (ID->WB)
 - Share the same fetch and control pipeline to save power (IF+cont.)
 - Similar to GPU's

30

ILP Techniques: Superscalar



31

Other Techniques for ILP: Out of Order Execution

- ❑ Eliminating stall conditions decreases CPI
- ❑ Reorder instructions to avoid stalls
- ❑ Example (5-stage LC2K pipeline):

add	1	2	3		add	1	2	3
lw	3	2	16		lw	3	2	16
nor	2	6	7		add	4	5	1
add	4	5	1		nor	2	6	7

Blue arrows indicate out-of-order execution: the second 'add' (4,5,1) is executed before the 'nor' (2,6,7) because its operands are available earlier.

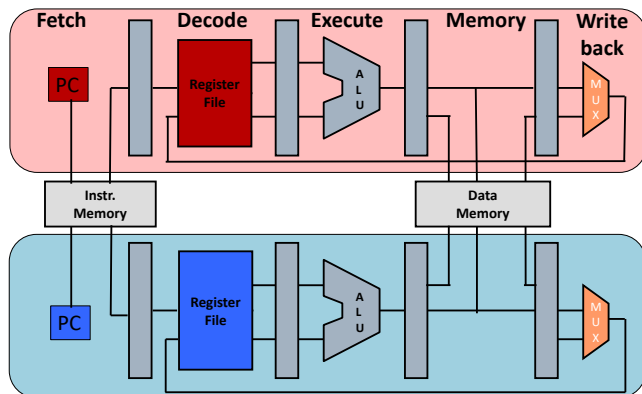
32

Why Use Out of Order Execution?

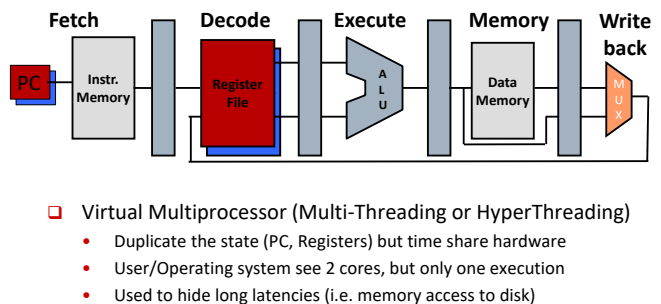
- ❑ Some instructions take a long time to execute
 - Floating point operations
 - Some loads and stores (more when we talk about memory hierarchy)
- ❑ Options:
 - Increase cycle time
 - Increase number of pipeline stages
 - Execute other instructions while you wait

33

TLP Techniques: Multiprocessors



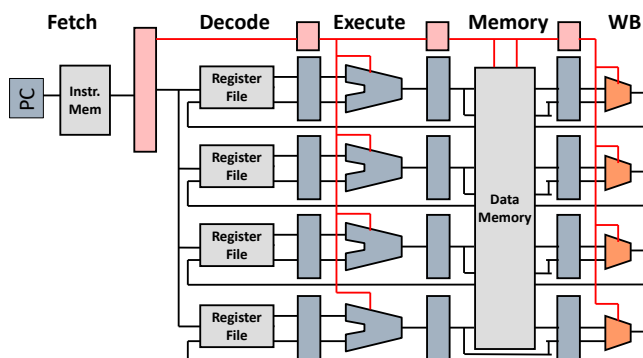
Other Techniques for TLP: Multi-Threading



- ❑ Virtual Multiprocessor (Multi-Threading or HyperThreading)
 - Duplicate the state (PC, Registers) but time share hardware
 - User/Operating system see 2 cores, but only one execution
 - Used to hide long latencies (i.e. memory access to disk)


35

DLP Techniques: Single Instr. Multiple Data (SIMD)



36

Building a GPU



- ❑ Add special functional units in EX
- ❑ Combine Techniques
 - SIMD + MP + MT = SIMT
- ❑ MT used to hide memory latencies
- ❑ SIMD used to decrease power of fetch/control
- ❑ MP used to improve throughput

37