


EECS 370

Virtual Memory Basics

Agenda

- Motivation
- Virtual Memory Principles
- Page Tables
- Class Problem

Storage Hierarchy



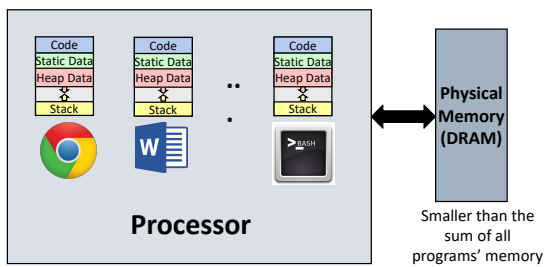
Memory: Two Issues

1. We've been working with the abstraction that all programs have full, private access to memory
 - But in practice, multiple programs run at the same time!

A screenshot of the Windows Task Manager 'Performance' tab, specifically the 'Memory' section. It shows a list of processes and their memory usage. The 'System' process is highlighted, showing it is using 47% of the 16 GB of installed memory.

- What happens if two programs try to write to the same memory address??

Revisit real system view—multitasking



Memory: Two Issues

2. Even if only one program is running, modern computers have 48-64 bit address spaces!
 - No computer actually has 18 exabytes (18 billion GBs)
 - What if a computer tries to write to address 0xFFFF...FFFF
 - Should it just crash??

Memory: Two Issues

- Modern systems use the same solution for both problems: **virtual memory**
 - In a nutshell, each program thinks it has full, private access to memory (it can safely index any address from 0x0-0xFFFF...FFFF)
 - Hardware and software transparently maps these addresses to distinct addresses in DRAM and in hard disk / SSD
 - Focus for the next 3 lectures

Agenda

- Motivation
- Virtual Memory Principles
- Page Tables
- Class Problem


Basics of Virtual Memory

- Any time you see the word **virtual** in computer science and architecture it means "using a level of indirection".


- Examples?



- Virtual memory hardware changes the virtual address the programmer sees into the physical one the memory chips see.



Virtual Memory




How to Translate Addresses?

- Address Translation is not done entirely in hardware
- We'll get help in software via the **operating system**
- The operating system is a special set of programs
 - Always running (after the system boots)
 - Is in charge of **managing the hardware resources** for all other running programs
 - E.g. initializing memory for a starting program, managing the file system, choosing when a particular program gets to run...
 - ... and translating virtual addresses into physical addresses!
- OS handles address translation by maintaining a data structure in main memory: the **page table**


Virtual memory terminology

- Memory is divided into fixed-size chunks called **Pages** (e.g., 4KB for x86)
 - Size of physical page = size of virtual page
- A virtual address consists of
 - A virtual page number
 - A page offset field (low order bits of the address)
- Translating a virtual address into a physical address only requires translating the page numbers
 - The page offset will stay the same



Page Table

- Translate page numbers using **page tables**
- Contains address translation information, i.e. virtual page # → physical page #
- Each process has its own page table
 - Maintained by operating system (OS)
- Page tables themselves are kept in memory by OS, and OS knows the physical address of the page tables
 - No address translation is required by the OS for accessing the page tables




Poll: What is the cost of this scheme? (select all that apply)

- a) Uses up more memory
- b) Takes longer to do loads and stores
- c) Fewer addresses accessible by program
- d) None of the above

Why Pages?

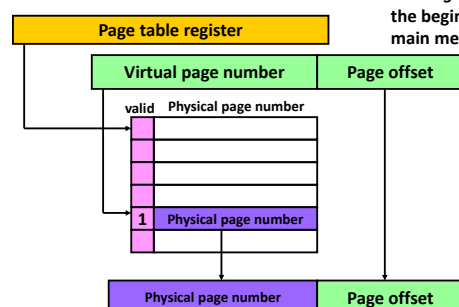
- Why have the idea of "pages"? Why not just have a mapping for each individual address?
 - Equivalent to asking: "why not have pages of size 1 byte?"
 - Otherwise - need a mapping entry for every single element of memory
 - The mapping data would take up as much space as the actual program data!
 - Also would screw up spatial locality of cache blocks (things contiguous in virtual memory wouldn't be contiguous in physical memory)



Agenda

- Motivation
- Virtual Memory Principles
- Page Tables**
- Class Problem

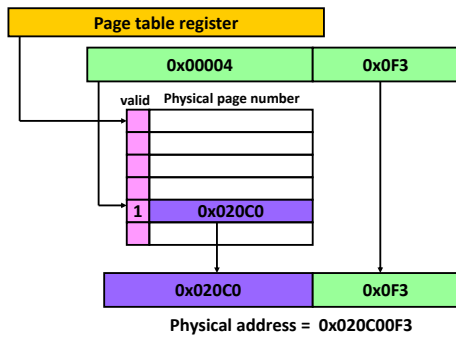
Page table components



The **Page table register** points to the beginning of the page table in main memory

Page table components - Example

Virtual address = 0x000040F3




Virtual Memory Goals

- VM should provide the following 3 capabilities to the programs:

- 1. Transparency**
 - Don't need to know how other programs are using memory
- 2. Protection**
 - No program can access the data of any other program
- 3. Programs not limited by DRAM capacity**
 - Each program can have more data than DRAM size

1-2: How to achieve transparency & protection?



3. How to be not limited by DRAM capacity?

- Use disk as temporary space in case memory capacity is exhausted
 - This temporary space in disk is called **swap partition** in Linux-based systems
 - For fun check swap space in a linux system by:

\$: top

```

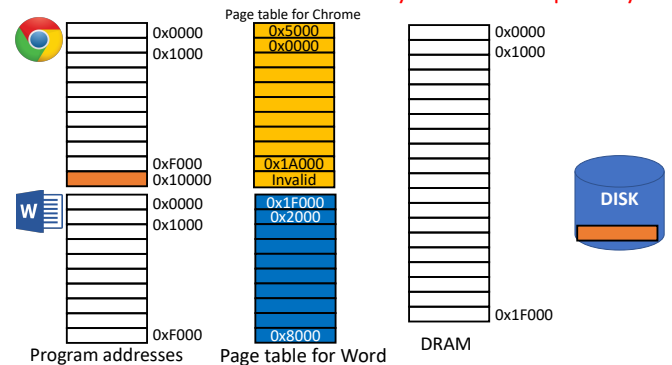
Tasks: 662 total, 1 running, 661 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.1 us, 0.0 sy, 0.0 ni, 99.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 32784372 total, 10813444 used, 21890928 free, 1018840 buffers
KiB Swap: 35162108 total, 89248 used, 35072860 free, 7053764 cached Mem

  PID USER      PR  NI  VIRT  RES  SHR S %CPU %MEM    TIME+  COMMAND
 60256 nehaag   20   0 25356 3356 2444 R  6.0  0.0   0:00.02 top
    1 root      20   0 38424  9040 2780 S  0.0  0.0   1:56.96 init
    2 root      20   0 0 0 0 S  0.0  0.0   0:02.21 kthreadd
    3 root      20   0 0 0 0 S  0.0  0.0   6:20.75 ksoftirqd/0
    
```

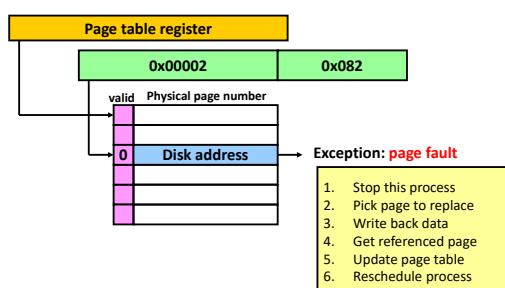
2. How to be not limited by DRAM capacity?

- We can mark a page table entry as "Invalid", indicating that the data for that page doesn't exist in main memory, but instead is located on the disk
- Looking up a page table entry that corresponds to disk is called a **page fault**

2. How to be not limited by DRAM capacity?



Page faults



How do we find it on disk?

- That is not a hardware problem! Go take EECS 482! ☺
- This is the operating system's job. Most operating systems partition the disk into logical devices (C:, D:, /home, etc.)
- They also have a hidden partition to support the disk portion of virtual memory
 - Swap partition** on UNIX machines
 - You then index into the correct page in the swap partition.

Agenda

- Motivation
- Virtual Memory Principles
- Page Tables
- **Class Problem**

Class Problem

- Given the following:
 - 4KB page size, physical memory of 16KB, page table stored in physical page 0 and can never be evicted, 20 bit, byte-addressable virtual address space.
 - The page table initially has virtual page 0 in physical page 1, virtual page 1 in physical page 2 and no valid data in other physical pages.
- Fill in the table on the next slide for each reference
 - Note: like caches we'll use LRU when we need to replace a page.



30



31

Class Problem (continued)

4KB → 12 bit → 3 hex offset 4 pages

virtual page table size: 20-12 = 8 bits = 256 (entries)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0	N	1 ~ (0 occupied)
0x01F0C	1	N	2 ~
0x20F0C	20 > 3	Y	3 ~
0x00100	0	N	1 ~
0x00200	0	N	1 ~
0x30000	30 > 3	Y	2 ~ (dirty, LRU)
0x01FFF	1	Y	3 ~ (dirty, LRU)
0x00200	0	N	1 ~

4KB page size, physical memory of 16KB, page table stored in physical page 0 and can never be evicted, 20 bit, byte-addressable virtual address space.
The page table initially has virtual page 0 in physical page 1, virtual page 1 in physical page 2 and no valid data in other physical pages.

Poll: How many hex digits should the page number be?

Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100	0x0	N	0x1100
0x00200	0x0	N	0x1200
0x30000	0x30	Y (into 2)	0x2000
0x01FFF	0x1	Y (into 3)	0x3FFF
0x00200	0x0	N	0x1200

4KB page size, physical memory of 16KB, page table stored in physical page 0 and can never be evicted, 20 bit, byte-addressable virtual address space.
The page table initially has virtual page 0 in physical page 1, virtual page 1 in physical page 2 and no valid data in other physical pages.



32



33

Size of the page table

- How big is a page table entry?
 - For 32-bit virtual address:
 - If the machine can support 1GB = 2^{30} bytes of physical memory and we use pages of size 4KB = 2^{12} ,
 - then the physical page number is 30-12 = 18 bits.
Plus another valid bit + other useful stuff (read only, dirty, etc.)
 - Let say about 3 bytes.
- How many entries in the page table?
 - 1 entry per virtual page
 - ARM virtual address is 32 bits – 12 bit page offset = 20
 - Total number of virtual pages = 2^{20}
- Total size of page table = Number of virtual pages
 - * Size of each page table entry = $2^{20} \times 3$ bytes ~ 3 MB



34