
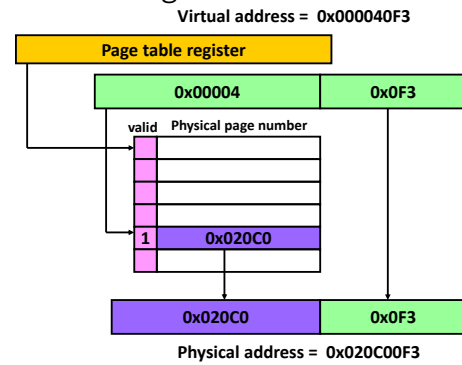


EECS 370

Multi-Level Page Tables



Reminder: Page tables



Agenda

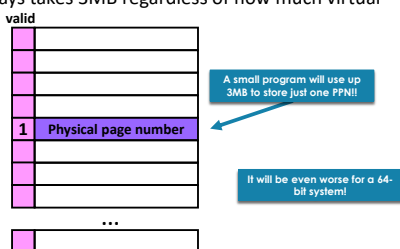
- Motivation for Multi-level Page Tables
- Example architecture
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design
- VM Miscellanea

Size of the page table

- How big is a page table entry?
 - For 32-bit virtual address:
 - If the machine can support $1\text{GB} = 2^{30}$ bytes of physical memory and we use pages of size 4KB = 2^{12} ,
 - then the physical page number is $30-12 = 18$ bits.
 - Plus another valid bit + other useful stuff (read only, dirty, etc.)
 - Let say about 3 bytes.
- How many entries in the page table?
 - 1 entry per virtual page
 - ARM virtual address is 32 bits – 12 bit page offset = 20
 - Total number of virtual pages = 2^{20}
- Total size of page table = Number of virtual pages
 - * Size of each page table entry
 - = $2^{20} \times 3 \text{ bytes} \sim 3 \text{ MB}$

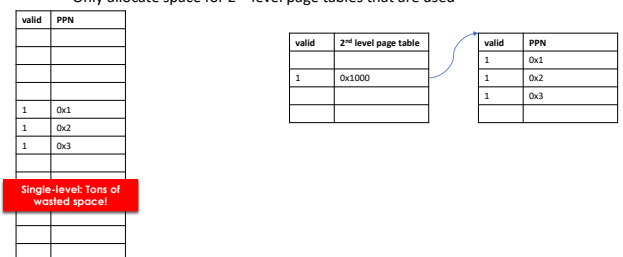
How can you organize the page table?

1. Single-level page table occupies continuous region in physical memory
 - Previous example always takes 3MB regardless of how much virtual memory is used



How can you organize the page table?

2. Option 2: Use a multi-level page table
 - 1st level page table (much smaller!) holds addresses 2nd level page tables
 - 2nd level page tables hold translation info, or 3rd level page tables if we wanna go deeper
 - Only allocate space for 2nd level page tables that are used



Multi-Level Page Table

- Only allocate second (and later) page tables when needed
- Program starts: everything is invalid, only first level is allocated

Single Level

valid	2 nd level page table
0	
0	
0	
0	

- As we access more, second level page tables are allocated

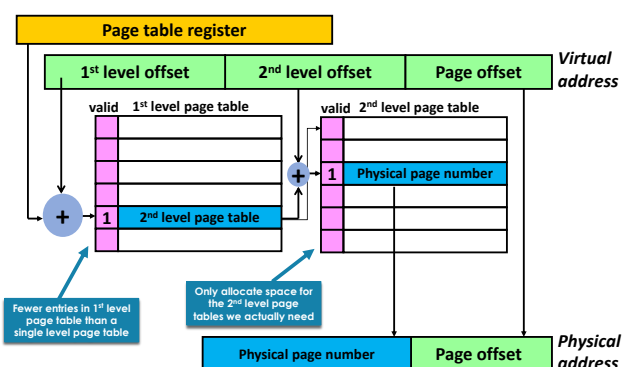


valid	2 nd level page table	valid	PPN	valid	PPN
1	0x1000	1	0x1	1	0x9a
1	0x3500	1	0x6	1	0x3
		1	0x2	0	
		1	0x1f	1	0xff

Multi-level: Size is proportional to amount of memory used

Common case: most programs use small portion of virtual memory space


Hierarchical page table




Agenda

- Motivation for Multi-level Page Tables
- **Example architecture**
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design
- VM Miscellanea

Hierarchical page table – 32bit Intel x86




- How many bits in the virtual 1st level offset field? 10
- How many bits in the virtual 2nd level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1st level page table? $2^{10}=1024$




11

Hierarchical page table – 32bit Intel x86




- Let's say physical address size + overhead bits is 4 bytes per entry
- Total size of 1st level page table
 - 4 bytes * 1024 entries = 4 KB



12

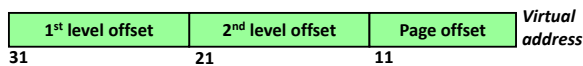
Hierarchical page table

- How many entries in the 2nd level of the page table?
 - $2^{10} = 1024$
- How many bytes for each VPN in a 2nd level table?
 - Let's round up to 4 bytes



13

Hierarchical page table – 32bit Intel x86



- How many bits in the virtual 1st level offset field? 10
- How many bits in the virtual 2nd level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1st level page table? $2^{10}=1024$
- How many bytes for each entry in the 1st level page table? 4
- How many entries in the 2nd level of the page table? $2^{10}=1024$
- How many bytes for each entry in a 2nd level table? ~4
- **What is the total size of the page table?** $4K+n*4K$
(here n is number of valid entries in the 1st level page table)

14

Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- **Class Problem: 32bit Intel x86**
- Class Problem: Multi-Level VM Design
- VM Miscellanea

Class Problem (32 bit x86)

- What is the least amount of memory that could be used? When would this happen?
- What is the most memory that could be used? When would this happen?
- How much memory is used for this memory access pattern:
0x00000ABC
0x00000ABD
0x10000ABC
0x20000ABC
- How much memory if we used a single-level page table with 4KB pages? Assume entries are rounded to the nearest word (4B)

16

Class Problem (32 bit x86)

- What is the least amount of memory that could be used? When would this happen?
 - 4KB for 1st level page table. Occurs when no memory has been accessed (before program runs)
- What is the most memory that could be used? When would this happen?
 - 4KB for 1st level page table
+ 1024*4KB for all possible 2nd level page tables
= 4100KB (which slightly greater than 4096KB)
 - Occurs when program uses all virtual pages (= 2^{20} pages)

17

Class Problem (32 bit x86)

- How much memory is used for this memory access pattern:

0x00000ABC // Page fault
 0x00000ABD
 0x10000ABC // Page fault
 0x20000ABC // Page fault

2^{12} , 3hex

- 4KB for 1st level page table + 3*4KB for each 2L page table
 = 16 KB

Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design**
- VM Miscellanea




18



19

Class Problem – Multi-level VM



- Design a two-level virtual memory system of a byte addressable processor with **24-bit long addresses**. No cache in the system. **256Kbytes of memory** installed, and no additional memory can be added.
 - Virtual memory page: 512 Bytes.** Each page table entry must be an integer number of bytes, and must be the smallest size required to fit the physical page number + 1 bit to mark valid-entry
 - We want each second-level page table to fit exactly in one memory page, and 1st level page table entries are 3 bytes each (a memory address pointer to a 2L page table).
- Compute:
 - Number of entries in each 2nd level page table;
 - Number of virtual address bits used to index the 2nd level page table;
 - Number of virtual address bits used to index the 1st level page table;
 - Size of the 1st level page table.

entry size: 10 bit
 round: 2B
 $512/2 = 256 \Rightarrow 8$
 $\Rightarrow 8 \text{ bit represent}$



20

Class Problem – Multi-level VM

Page Offset: 9 bits (512B page size)
 Physical address = 18b (256KB Mem size)
 Physical page number = 18b (256KB mem size) – 9b (offset)

Physical page number = 9b Page offset = 9b

2nd level page table entry size: 9b (physical page number) + 1b = ~2 bytes
 2nd level page table fits exactly in 1 page
 #entries in 2nd level page table is 512 bytes / 2 bytes = 256
 #entries in 2nd level page table = 256 \rightarrow Virtual page bits = 8b

Virtual 1st level page bits = 24 – 8 – 9 = 7b
 1st level page table size = $2^7 * 3 \text{ bytes} = 384B$

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b



22

Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design
- VM Miscellanea**

Page Replacement Strategies

- Page table indirection enables a fully associative mapping between virtual and physical pages.
- How do we implement LRU in OS?
 - True LRU is expensive, but LRU is a heuristic anyway, so approximating LRU is fine
 - Keep a “**accessed**” bit per page, cleared occasionally by the operating system. Then pick any “unaccessed” page to evict



23



24

Other VM Translation Functions

- Page data location
 - Physical memory, disk, uninitialized data
- Access permissions
 - Read only pages for instructions
 - This is how your system detects segmentation faults**
- Gathering access information
 - Identifying dirty pages by tracking stores
 - Identifying accesses to help determine LRU candidate



25